

Exceptional service in the national interest

FROM EPETRA TO TPETRA

Migrating Your Code with Trilinos

Christopher Siefert, Carl Pearson and Timothy Fuller

EuroTUG '24



WHY THE TPETRA TRANSITION?

WHY THE TPETRA TRANSITION?

- Epetra does not...
 - Work on GPU platforms (or work well with OpenMP).
 - Handle over 2.1B unknowns.
 - Allow for non-double Scalar types (e.g. float, complex, AD types).
 - Make optimal use of vendor TPLs.
- Tpetra does all of these (and more).
- Epetra stack to be removed from Trilinos in FY25.



TPETRA STACK IS TESTED NIGHTLY ON A LOT OF PLATFORMS

CPU

Amber (Sapphire Rapids/DDR)
Eclipse (Broadwell)
Rocinante (Sapphire Rapids/HBM)
Stria (Arm Cavium Thunder-X2)

AMD GPU

Frontier (MI250) Tioga (MI250) RZVernal (MI300)

NVIDIA GPU

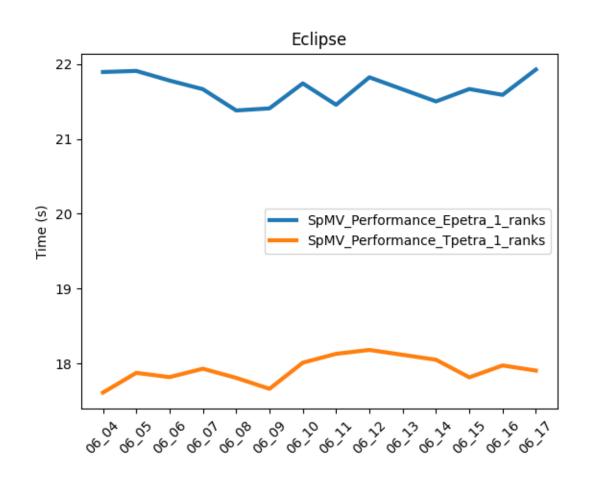
Vortex (V100) Perlmutter (A100) Hops (H100)

Intel GPU

Sunspot (Intel Max) coming soon



EVEN ON CPUS, TPETRA IS FASTER





Tpetra is roughly 20% faster for SpMV!

EPETRA TO TPETRA WITHOUT KOKKOS

CASE STUDY: SNL'S ALEGRA CODE

- Goal
 - Move from Epetra/AztecOO/ML to Tpetra/Belos/MueLu.
 - Enable > 2.1 billion unknowns.
 - Not interested in GPU porting.
- Alegra's Strategy
 - Support both Epetra and Tpetra via runtime switch.
 - Compare results w/ nightly testing.
 - Work with Trilinos developers to develop missing capabilities.



EPETRA-TO-TPETRA: MAP

Epetra_Map * map = new Epetra_Map(-1,num_owned, &(indices.front()), 0, Comm);

```
using LO = Tpetra::Map<>::local_ordinal_type;
using GO = Tpetra::Map<>::global_ordinal_type;
using NT = Tpetra::Map<>::node_type;
const Tpetra::global_size_t invalid = Teuchos::OrdinalTraits<Tpetra::global_size_t>::invalid();
using map_type = Tpetra::Map<LO,GO,NT>;

auto map = Teuchos::rcp(new map_type(invalid,indices(),0,comm));
auto map_owned_plus_shared = Teuchos::rcp(new map_type(invalid,indices_owned_plus_shared(),0,comm));
```



EPETRA-TO-TPETRA: FECRSGRAPH

```
Epetra_FECrsGraph * graph = new Epetra_FECrsGraph(Copy,*map,&nnz.front());
for (auto item: items) {
  graph->InsertGlobalIndices(1,&grid, gcids.size(), &gcids.front());
graph->GlobalAssemble(); graph->OptimizeStorage();
using graph_type = Tpetra::FECrsGraph<LO,GO,NT>;
auto graph = Teuchos::rcp(new graph_type(map,map_owned_plus_shared,MAX_NNZ_ROW));
Tpetra::beginAssembly(*graph);
for (auto item: items) {
  graph->insertGlobalIndices(grid, gcids());
Tpetra::endAssembly(*graph);
```



EPETRA-TO-TPETRA: FEVECTOR

```
using my_type = Tpetra::FEMultiVector<SC,LO,GO,NT>;
auto vec = Teuchos::rcp(new mv_type(map,graph->getImporter(),1));
size_t len = vec->getMap()->getLocalNumElements();
auto data = vec->get1dViewNonConst();

for (int i=0; i<len; i++)
    data[i] = <whatever>;
```



EPETRA-TO-TPETRA: FECRSMATRIX

```
Epetra_FECrsMatrix * matrix = new Epetra_FECrsMatrix(Copy,*graph);
for (auto item: items) {
  matrix->ExtractGlobalRowView(grid,num_entries,values,indices);
  for(int j=0; j<num_entries; j++) values[j] = <whatever>;
matrix->OptimizeStorage();
using matrix_type = Tpetra::FECrsMatrix<SC,LO,GO,NT>;
auto matrix = Teuchos::rcp(new matrix_type(graph));
Tpetra::beginAssembly(*matrix);
for (auto item: items) {
  matrix->getLocalRowView(grid, indices, values);
  for(int j=0; j<indices.size(); j++) values[j] = <whatever>;
Tpetra::endAssembly(*matrix);
```



LESSONS LEARNED

- For performance reasons, Tpetra requires you do more work than Epetra does
 - Max nnz per row is required for graphs (avoid in-loop dynamic memory allocation)
 - FECrsGraph: owned_plus_shared map avoid link-list structures.
- Use typedefs/using: You can wind up with looooooong code lines otherwise.
- Alegra team worked with Trilinos devs to improve user interfaces and add needed features to Belos & MueLu.
 - Alegra was the one of the first users of the Tpetra::FE* tools, so they helped design better interfaces.
 - Some new features were "compatibility modes" in the new stack, others were algorithms which hadn't yet been ported to Tpetra.



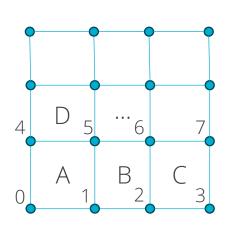


TPETRA'S FINITE ELEMENT ASSEMBLY EXAMPLE

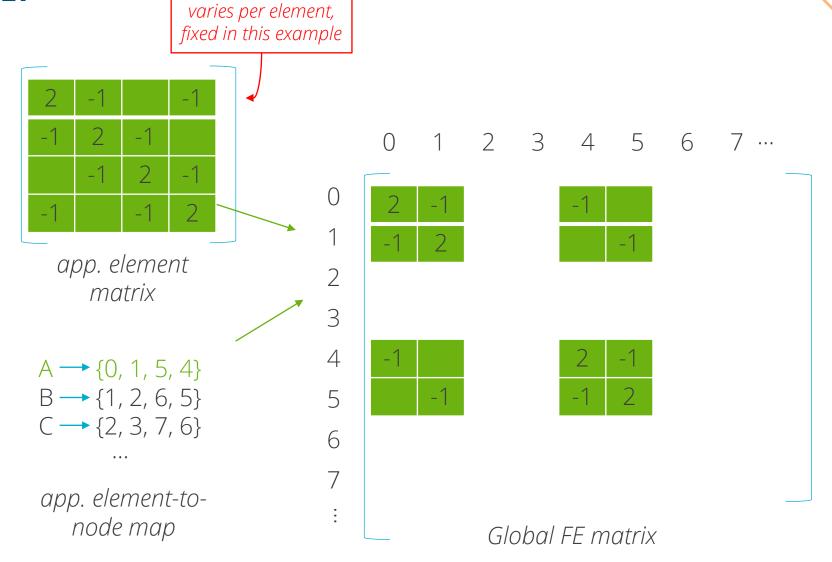
- Background on Kokkos
 - Kokkos Lectures & Slides (https://kokkos.github.io/kokkos-core-wiki/videolectures.html)
 - Single code for all CPUs, GPUs, and whatever else Kokkos supports.
- Example: <u>Trilinos/packages/tpetra/core/example/Finite-Element-Assembly</u>
- Application provides (our example mocks these)
 - Map of elements to nodes in global indices.
 - Methods for computing element matrices.
- Type-1 assembly
 - Local elements contribute to off-rank FE matrix rows for off-rank nodes.
 - Does not require ghosting of element state.
- No worksetting



FE ASSEMBLY IN BRIEF

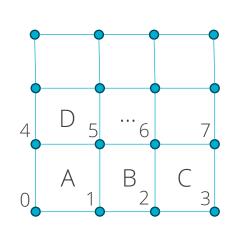


app. discretization

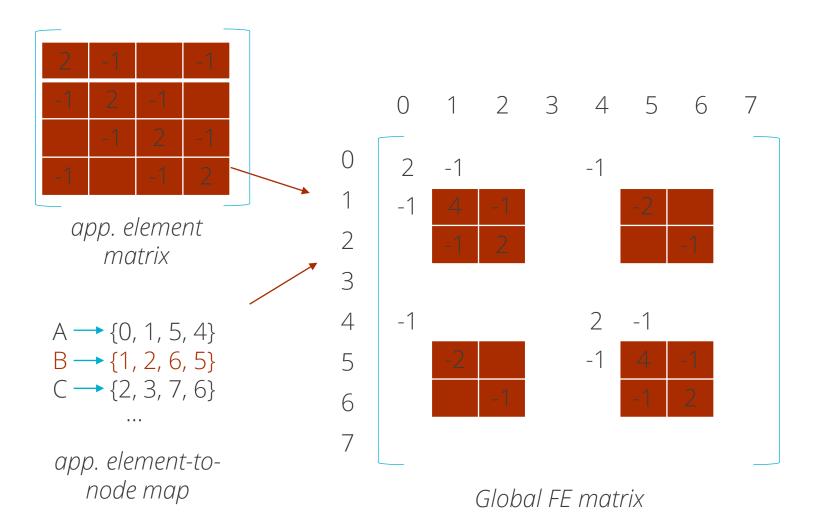




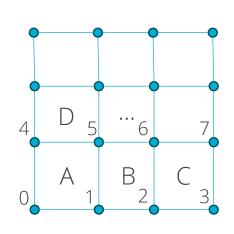
ELEMENT B'S CONTRIBUTION



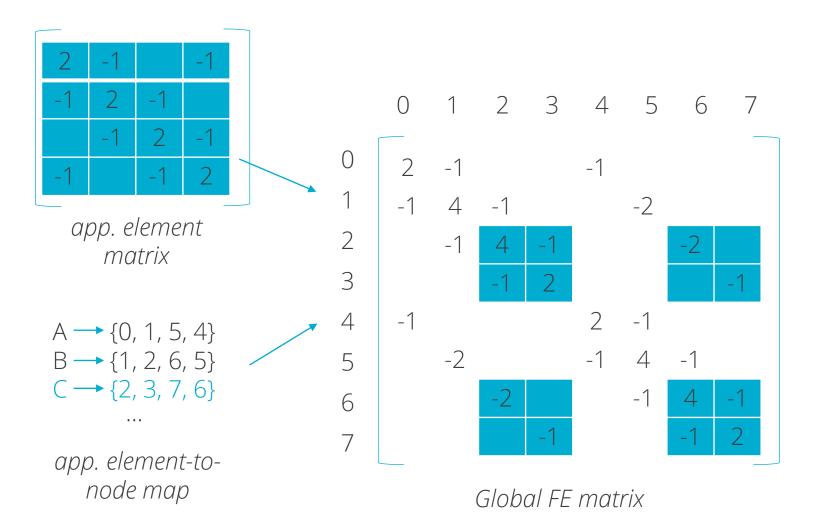
app. discretization



ELEMENT C'S CONTRIBUTION



app. discretization





FIVE CHANGES TO WATCH OUT FOR

Bare for-loops



- Kokkos::parallel_for
 - Allow device execution
 - Supports CPU execution too

Host allocations



Kokkos::View

Functions



KOKKOS_FUNCTION annotation

- Tpetra::[...]::getHostView()
 - Convenient use of global indices



- Tpetra::[]::getDeviceView()
 - Easier to use local indices "on device"

Normal addition



Atomic addition



HOST LOOP -> KOKKOS::PARALLEL_FOR

```
Kokkos::View<local_ordinal_type[4][4], hostType>
Teuchos::Array<Scalar> element_rhs(4);
Teuchos::Array<global_ordinal_type>
Tnetra::heginAssembly(*fe matrix *rhs):
  ReferenceQuad4(element matrix);
  ReferenceQuad4RHS(element_rhs);
         element node idx < nodesPerElem;</pre>
      owned_element_to_node_gids(
  for (size_t element_node_idx = 0;
```

Tpetra: a single active thread, loops over each local element

Kokkos: operate on local elements in parallel

Also works on host, singlethreaded (Kokkos::Serial) or multi-threaded

(::OpenMP, ::Threads)

```
Tpetra::Access::OverwriteAll);
auto localMatrix = fe matrix->getLocalMatrixDevice();
auto all element rhs unmanaged =
 makeUnmanaged(all element rhs);
 makeUnmanaged(all_element_matrix);
auto all lcids unmanaged = makeUnmanaged(all lcids);
Kokkos::parallel for
  Kokkos::RangePolicy<execution space, int>(
    0, numOwnedElements),
  KOKKOS LAMBDA (const size t element idx) {
 const pair type location pair(
   nodesPerElem*element_idx,
   nodesPerElem*(element idx+1));
 auto element matrix = Kokkos::subview(
   Kokkos::ALL);
 auto element_lcids = Kokkos::subview(
   all lcids unmanaged, location pair);
 auto element_rhs = Kokkos::subview(
   all_element_rhs_unmanaged, location_pair);
 ReferenceOuad4(element matrix):
 ReferenceQuad4RHS(element rhs);
```

Tpetra + Kokkos



HOST ALLOCATIONS -> KOKKOS VIEW

```
Kokkos::View<local_ordinal_type[4][4], hostType>
Teuchos::Array<Scalar> element_rhs(4);
Teuchos::Array<global ordinal type>
 Feuchos::Array<Scalar> column_scalar_values(4);
Tpetra::beginAssembly(*fe_matrix,*rhs);
     element_gidx < numOwnedElements;</pre>
  ReferenceQuad4(element_matrix);
  ReferenceQuad4RHS(element rhs);
         element_node_idx < nodesPerElem;</pre>
      owned element to node gids(
  for (size t element node idx = 0;
       element node idx < 4;
       ++element_node_idx) {
```

Tpetra: allocate some scratch space on the stack. Reused for each iteration of element loop

Kokkos: allocate enough device memory for all active threads

Kokkos: each thread gets its own piece of the preallocated scratch space

```
scalar_2d_array_type all_element_matrix(
   "all_element_matrix",nodesPerElem*numOwned:lements);
scalar_1d_array_type all_element_rhs(
   "all_element_rhs",nodesPerElem*numOwnedElements);
   "all_element_rhs",nodesPerElem*numOwnedElements);

Tpetra::beginAssembly(*fe_matrix,*rhs);
auto owned_element_to_node_gids =
   mesh.getOwnedElementToNode().getDeviceView(
```

```
KOKKOS_LAMBDA (const size_t element_idx) {
const pair_type location_pair(
    nodesPerElem*element_idx,
    nodesPerElem*(element_idx+1));

auto element_matrix = Kokkos::subview(
    all_element_matrix_unmanaged, location_pair,
    Kokkos::ALL);
auto element_lcids = Kokkos::subview(
    all_lcids_unmanaged, location_pair);
auto element_rhs = Kokkos::subview(
    all_element_rhs = Kokkos::subview(
    all_element_rhs_unmanaged, location_pair);

ReferenceQuad4(element_matrix);
ReferenceQuad4RHS(element_rhs);

for (int element_node_idx = 0;
    alement_node_idx = 0;
    alement_node_idx = node_ParElement_node_idx = 0;
    alement_node_idx = node_ParElement_node_idx
```

Tpetra + Kokkos



HOST FUNCTIONS -> KOKKOS_FUNCTION

```
Kokkos::View<local_ordinal_type[4][4], hostType>
Teuchos::Array<Scalar> element_rhs(4);
Teuchos::Array<global_ordinal_type>
Teuchos::Array<Scalar> column scalar values(!);
Tpetra::beginAssembly(*fe_matrix,*rhs);
  ReferenceQuad4(element_matrix)
  ReferenceQuad4RHS(element_rhs):
         element node idx < nodesPerElem;</pre>
    column_global_ids[element_node_idx] =
      owned_element_to_node_gids(
  for (size_t element_node_idx = 0;
       element_node_idx < 4;</pre>
```

Tpetra: fill scratch space with matrix for this element

Kokkos: each thread fills scratch space in parallel

These functions must be allowed to execute on the device

```
KOKKOS_FUNCTION
void Reference4Quad(...) {
   ...
}:
```

```
nodesPerElem*element idx,
  nodesPerElem*(element idx+1));
auto element matrix = Kokkos::subview(
  all_element_matrix_unmanaged, location_pair,
auto element lcids = Kokkos::subview(
  all_lcids_unmanaged, location_pair);
auto element_rhs = Kokkos::subview(
  all_element_rhs_unmanaged, location_pair);
ReferenceQuad4(element_matrix);
ReferenceQuad4RHS(element rhs);
     element node_idx < nodesPerElem;</pre>
      owned element to node gids(
for (int element node idx = 0;
     element node idx < nodesPerElem;</pre>
```

Tpetra Tpetra + Kokkos



GLOBAL INDICES -> LOCAL INDICES

```
element_gidx < numOwnedElements;</pre>
  ++element_gidx) {
ReferenceQuad4(element matrix);
ReferenceQuad4RHS(element rhs);
       element node idx < nodesPerElem;</pre>
  column global ids[element node idx] =
    owned element to node gids(
  global_ordinal_type global_row_id =
    owned_element_to_node_gids(
  fe_matrix->sumIntoGlobalValues(
    global_row_id, column_global_ids,
    column scalar values);
```

Tpetra: interact with FE matrix and RHS through global indices. Simpler interface, made possible by dynamic memory allocation

Kokkos: need device views of the local FE matrix and RHS to operate on.

Kokkos: local FE matrix and RHS only understands local indices.
Use Tpetra::Maps to translate between local and global

```
Tpetra::Access::ReadOnly);
rhs->getLocalViewDevice(
auto localMatrix = fe matrix->getLocalMatrixDevice();
auto all_element_rns_unmanaged =
 makeUnmanaged(all element rhs);
ReferenceQuad4(element matrix);
 ReferenceQuad4RHS(element rhs);
      element_node_idx < nodesPerElem;</pre>
      element_node_idx < nodesPerElem;</pre>
   const local_ordinal_type local_row_id =
     localMap.getLocalElement(owned_element_to_node_gits()
   for (int col_idx = 0; col_idx < nodesPerElem;</pre>
     localMatrix.sumIntoValues(local_row_id,
```

Tpetra + Kokkos



ATOMIC ADDITION

```
column_global_ids[element_node_idx] =
 owned_element_to_node_gids(
global_ordinal_type global_row_id =
 owned_element_to_node_gids(
   element matrix(element node idx, col idx):
fe matrix->sumIntoGlobalValues(
rhs->sumIntoGlobalValue(
```

Tpetra: contribute element values to FE matrix and RHS

Kokkos: atomic adds, since each thread (element) may contribute to the same node at the same time

```
element node idx < nodesPerElem:</pre>
   ++element node idx) {
  localColMap.getLocalElement(
    owned_element_to_node_gids(
   element node idx < nodesPerElem;</pre>
const local_ordinal_type local_row_id =
  localMap.getLocalElement(owned element to node gids(
for (int col idx = 0; col idx < nodesPerElem;</pre>
  localMatrix.sumIntoValues(local row id,
                             &(element_matrix(
Kokkos::atomic add(
  &(localRHS(local_row_id, 0)),
```

Tpetra + Kokkos



WRAP UP: ADDING KOKKOS TO TPETRA

- Create parallel execution using Kokkos::parallel_for
 - Also supports host CPU execution
- Kokkos::View for data accessed in parallel regions
 - Convert std::vector, Teuchos::Array, malloc, new, ...
- Functions called in that region must be KOKKOS_FUNCTION
 - e.g. producing the element matrix
 - ...and any data it requires must be in a Kokkos::View (material properties, node coordinates, etc.)
- Use Tpetra::[]::getDeviceView() to get Kokkos::View of Tpetra data
 - As a consequence, have to operate with local rather than global indices
- Parallel regions may require atomics for their contributions





CONCLUSIONS

- Moving to Tpetra has a lot going for it...
 - GPU portability.
 - > 2.1B unknowns.
 - Better support for complex, AD types.
 - Tested on many architectures.
 - Faster SpMVs on CPU platforms.
- It will require more work (especially with graph assembly), but you gain in performance.
- Can be used with or without explicit use of Kokkos.
 - But if you want GPU support, you'll need to get familiar with Kokkos.

